

Computer Vision - 2026


Week #06. Segmentation and Object Detection

Lectures by Alexei Kornaev ^{1,2}
Practical sessions by Kirill Yakovlev ^{1,3}

¹The Center for Top-Level Educational Programs
in AI, Innopolis University (IU), Innopolis

²RC for AI, National RC for Oncology
n.a. N.N.Blokhin, Moscow

³Ak Bars Bank, Kazan



Agenda

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

- 1 Outcomes
- 2 Introduction
- 3 Semantic Segmentation
- 4 Object Detection
- 5 Segment Anything Model (SAM)
- 6 Conclusion

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Section 1. Outcomes

Outcomes

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

By the end of this week, students will be able to:

- 1 **Formulate** segmentation and detection as structured prediction problems with precise mathematical notation.
- 2 **Explain** semantic segmentation architectures (U-Net, DeepLab) and the YOLO detection pipeline (backboneâneckâhead, loss, NMS).
- 3 **Implement** a YOLO detector and a semantic segmentation model, including data preparation, loss selection, and evaluation.
- 4 **Interpret metrics** (mIoU, Dice, etc.) and diagnose common failure modes.
- 5 **Select appropriate models** for real-time systems, robotics, medical imaging, or high-accuracy offline pipelines.

Key Takeaway

Segmentation and detection are structured vision problems requiring joint optimization of classification and spatial localization; model choice is a trade-off between accuracy, latency, and supervision cost.

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Section 2. Introduction

Types of Image Segmentation

CV-2026

A.Korinaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Segmentation Overview

Image segmentation is the process of partitioning an image into meaningful regions. It plays a key role in medical imaging, autonomous driving, and scene understanding.

- **Semantic Segmentation:** Classifies each pixel into predefined object categories.
- **Instance Segmentation:** Identifies individual object instances within categories.
- **Panoptic Segmentation:** Unifies semantic and instance segmentation by labeling both things (objects) and stuff (background regions).





Figure: Types of segmentation problems.

Problem Formulation: Structured Outputs



Input $x \in \mathbb{R}^{H \times W \times 3}$ (H, W : res., K : classes, θ : params, $u = (i, j)$: pixel, N : objects).


Semantic Segmentation

$$f_{\theta}(x) = P \in [0, 1]^{H \times W \times K}, P(u, k) = \mathbb{P}(y_u = k \mid x).$$

Output $\hat{Y} \in \{1, \dots, K\}^{H \times W}$ via $\hat{y}_u = \arg \max_k P(u, k)$.

Object Detection

$$f_{\theta}(x) = \{(b_i, c_i, s_i)\}_{i=1}^N, \text{ box } b_i \in \mathbb{R}^4, \text{ class } c_i, \text{ score } s_i \in [0, 1].$$



Instance Segmentation

$$f_{\theta}(x) = \{(m_i, c_i, s_i)\}_{i=1}^N, \text{ mask } m_i \in \{0, 1\}^{H \times W}, \text{ class } c_i, \text{ score } s_i.$$

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Taxonomy: Tasks and Outputs

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Four Canonical Tasks

Task	Output	Supervision	Typical Models
Semantic Segm.	Dense per-pixel class map $H \times W$	Pixel labels (stuff + things)	FCN, <u>U-Net</u> , DeepLab, SegFormer
Instance Segm.	Set of masks $\{m_i\}$ + classes	Masks per instance (expensive)	Mask R-CNN, SOLO, Mask2Former
Panoptic Segm.	Single map with instance IDs + stuff	Stuff + instance masks (hardest)	Panoptic-FPN, Mask2Former, DETR
Object Detection	Set of boxes $\{b_i\}$ + classes	Boxes per instance (cheaper)	Faster R-CNN, <u>YOLO</u> , RetinaNet, DETR

Implementation View: Data and Annotation Pipelines

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Data Shapes

- **Detection dataset:** $\{(x^{(n)}, \{(b_i^{(n)}, c_i^{(n)})\}_{i=1}^{N_n})\}_{n=1}^M$
- **Segmentation dataset:** $\{(x^{(n)}, y^{(n)}), y^{(n)} \in \{1, \dots, K\}^{H \times W}\}_{n=1}^M$
- **Instance dataset:** $\{(x^{(n)}, \{(m_i^{(n)}, c_i^{(n)})\})\}_{n=1}^M$

Annotation Cost (Engineering Reality)

- Boxes: fast labeling, scalable to large datasets (COCO-like).
- Masks: slow labeling, often smaller datasets / domain-specific (medical, industrial QC).
- Panoptic: most expensive; usually limited data, heavy augmentation.

Losses: What We Optimize (Overview)

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Segmentation (dense)

- **Cross-Entropy (CE):** pixel-wise classification.
- **Dice / IoU losses:** robust for class imbalance and small structures.

$$\mathcal{L}_{seg} = \underbrace{\lambda_{CE}}_{\text{green}} \mathcal{L}_{CE} + \underbrace{\lambda_{Dice}}_{\text{green}} \mathcal{L}_{Dice}$$

Detection (set of objects)

- Classification loss: CE / focal loss (handles imbalance: many negatives).
- Box regression: SmoothL1, IoU/GIoU/DIoU/CIoU family.

$$\mathcal{L}_{det} = \mathcal{L}_{cls} + \underbrace{\lambda_{box}}_{\text{green}} \mathcal{L}_{box} + \underbrace{\lambda_{obj}}_{\text{green}} \mathcal{L}_{obj}$$

Metrics: What We Report (and Why It Matters)

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Segmentation Metrics

$$IoU = \frac{|Y \cap \hat{Y}|}{|Y \cup \hat{Y}|}, \quad Dice = \frac{2|Y \cap \hat{Y}|}{|Y| + |\hat{Y}|}$$

- **mIoU**: mean IoU over classes (common in semantic segmentation).
- **Dice**: widely used in medical imaging (small targets).

Detection Metrics (COCO-style)

- **AP@0.5**: easier metric (Pascal VOC tradition).
- **AP@[0.5:0.95]**: strict metric; averages over IoU thresholds.
- **AR**: average recall; highlights missed objects.

Analytics: Typical Failure Modes by Task

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Structured Error Taxonomy

- **Segmentation:** boundary errors, small structures disappear, class imbalance, label noise, domain shift.
- **Detection:** localization error, confusion between similar classes, duplicates, missed small/occluded objects.
- **Instance/Panoptic:** overlapping instances, crowd scenes, thin objects, inconsistent instance IDs.

Symptom	Metric signature	Likely cause / fix
Miss small objects	low AR, AP_small	increase resolution, FPN, augment, re-weight samples
Too many duplicates Blurry boundaries	AP drops, many FP mIoU ok, Dice low	tune NMS / use end-to-end set models add boundary loss / higher-res decoder / better aug
Domain shift	all metrics drop	adaptation, better pretraining, prompt-able models (SAM)

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Section 3. Semantic Segmentation

Semantic Segmentation: Core Idea

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Goal

Predict a **dense** label map: each pixel $u = (i, j)$ gets a semantic class.

$$y \in \{1, \dots, K\}^{H \times W}, \quad f_{\theta}(x) = P \in [0, 1]^{H \times W \times K}$$

$$P(u, k) = \mathbb{P}(y_u = k \mid x), \quad \hat{y}_u = \arg \max_k P(u, k)$$

What makes it non-trivial?

- Need **global context** (class identity) + **local precision** (boundaries).
- Strong **class imbalance** (background dominates), small objects/structures.
- Output is **structured**: neighboring pixels are correlated.

Typical families

Encoder-decoder (U-Net), atrous/ASPP (DeepLab), transformer-based dense heads (SegFormer).

U-Net Architecture

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

- U-Net consists of a contracting path (encoder) and an expansive path (decoder).
- Skip connections concatenate corresponding encoder and decoder feature maps to retain spatial details.
- Final layer applies a 1×1 convolution to obtain pixel-wise classification.




Figure: Detailed U-Net structure [Ronneberger et al., 2015].

Residual Connections in U-Net

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Key Idea (He et al., 2015)

Residual connections help mitigate the vanishing gradient problem and improve gradient flow in deep networks.

- A residual connection allows the gradient to bypass some layers, making training easier.
- Instead of learning the full transformation $F(x)$, the network learns the residual function $F(x) + x$.
- Some U-Net variants, such as ResUNet, integrate residual blocks for better performance.




Figure: Residual connections in a CNN block.

Upsampling in U-Net

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

- U-Net uses upsampling to restore spatial resolution in the decoder.
- Two main types of upsampling:
 - **Interpolation-based Upsampling:** Bilinear or nearest neighbor interpolation followed by convolution.
 - **Transposed Convolution (Deconvolution):** Learnable filters upscale feature maps.




Figure: Nearest neighbor interpolation [Prince, 2023].

Transposed Convolution (Deconvolution)

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

- Transposed convolution is a learnable upsampling method.
- Instead of direct interpolation, it applies convolution on an expanded input.
- Steps:
 - ① Insert zeros between pixels (if stride $\neq 1$).
 - ② Convolve with a learnable kernel.
 - ③ Output has a larger spatial resolution.




Figure: Example of transposed convolution operation.

U-Net: Encoder–Decoder with Skip Connections

CV-2026

A.Korinaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Main idea (U-Net)

Contracting path learns semantics; **expanding path** restores resolution. **Skip connections** fuse low-level spatial detail with high-level context.

Computation pattern (abstract)

Let E_ℓ be encoder features, D_ℓ decoder features:

$$E_{\ell+1} = \phi(\text{Down}(E_\ell)), \quad D_\ell = \phi(\text{Up}(D_{\ell+1}) \oplus E_\ell)$$

where $\oplus$ is concatenation (or sum), ϕ is Conv+Norm+Act. Final prediction:
 $P = \text{Softmax}(\text{Conv}_{1 \times 1}(D_0))$.

Implementation notes

- Up-sampling: bilinear+conv (stable) or transposed conv (learnable).
- Normalization: BatchNorm / GroupNorm (small batch) for segmentation.
- Deep supervision sometimes helps: auxiliary heads at intermediate scales.

U-Net: Training Algorithm

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Algorithm 1 U-Net Training Procedure

Require: Training dataset $\mathcal{D} = \{(x^{(n)}, y^{(n)})\}_{n=1}^M$

Require: Parameters θ , learning rate η

```
1: for epoch = 1, ..., E do
2:   for mini-batch  $\{(x_b, y_b)\}_{b=1}^B$  do
3:      $P_b \leftarrow f_{\theta}(x_b)$  ▷ Forward pass (encoder-decoder)
4:      $\mathcal{L}_{CE} \leftarrow \text{CrossEntropy}(P_b, y_b)$ 
5:      $\mathcal{L}_{Dice} \leftarrow \text{DiceLoss}(P_b, y_b)$ 
6:      $\mathcal{L} \leftarrow \lambda_{CE} \mathcal{L}_{CE} + \lambda_{Dice} \mathcal{L}_{Dice}$ 
7:      $\theta \leftarrow \theta - \eta \nabla_{\theta} \mathcal{L}$  ▷ Backpropagation
8:   end for
9: end for
10: return trained parameters  $\theta$ 
```

Remarks

- Encoder may be pretrained (ImageNet initialization).
- Data augmentation applied jointly to (x, y) before forward pass.

DeepLab: Atrous Convolution and ASPP

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Main idea (DeepLab family)

Increase receptive field **without downsampling too much** using atrous (dilated) conv, then aggregate multi-scale context with **ASPP**.

Atrous (dilated) convolution

For 1D intuition (extends to 2D):

$$y[t] = \sum_k w[k] x[t + r \cdot k]$$

r is dilation rate ($r = 1$ is standard conv). Larger r increases receptive field.

ASPP (multi-scale context)

Parallel atrous conv branches with different dilation rates $\{r_1, r_2, r_3, \dots\}$:

$$\text{ASPP}(F) = \text{Concat}(\text{Conv}_{r_1}(F), \text{Conv}_{r_2}(F), \dots)$$

Training Setup: Data, Augmentations, and Splits

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Data formats

- Input: image $x \in \mathbb{R}^{H \times W \times 3}$
- Label map: $y \in \{1, \dots, K\}^{H \times W}$ (or ignore label 0 / 255)

Batching strategy (practical)

- Random crop to fixed size (e.g., 512×512), keep aspect via resize+pad.
- Balance rare classes: class-aware sampling or oversample patches containing target.

Augmentations (segmentation-safe)

- Geometric: flips, rotations, scale jitter, random crop (apply to image and mask).
- Photometric: brightness/contrast, blur/noise (image only).
- CutMix/MixUp variants exist but need care with masks (often helps with domain shift).

Losses for Semantic Segmentation

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Pixel-wise Cross-Entropy (multiclass)

Let $p_u(k) = P(u, k)$ and one-hot target $\mathbf{1}[y_u = k]$: $\mathcal{L}_{CE} = -\frac{1}{HW} \sum_u \sum_{k=1}^K \mathbf{1}[y_u = k] \log p_u(k)$

Class-weighted CE for imbalance: $\mathcal{L}_{WCE} = -\frac{1}{HW} \sum_u w_{y_u} \log p_u(y_u)$

Dice loss (binary; per-class for multiclass)

For a class k (soft prediction $\hat{y}_u = p_u(k)$, target $y_u \in \{0, 1\}$):

$$Dice_k = \frac{2 \sum_u y_u \hat{y}_u + \epsilon}{\sum_u y_u + \sum_u \hat{y}_u + \epsilon}, \quad \mathcal{L}_{Dice} = 1 - \frac{1}{K} \sum_{k=1}^K Dice_k$$

Common composite objective

$$\mathcal{L} = \lambda_{CE} \mathcal{L}_{(W)CE} + \lambda_{Dice} \mathcal{L}_{Dice}$$

Optimization Recipe (Baseline That Works)

CV-2026

A.Korbaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Baseline training loop (high-level)

- Backbone pretraining: ImageNet (common) or domain-specific (medical).
- Optimizer: AdamW (transformers) or SGD+momentum (CNNs).
- LR schedule: poly decay / cosine; warmup for transformers.
- Regularization: weight decay; dropout mainly in transformer decoders.

Typical hyperparameters (starting point)

- Crop: 512^2 (Cityscapes) / 256^2 – 512^2 (medical, depends on resolution)
- Batch: as large as fits; if small, prefer GroupNorm and gradient accumulation
- $\lambda_{CE} = 1$, $\lambda_{Dice} \in [0.5, 2]$ (tune per dataset imbalance)

Two common failure modes

- (1) Over-smoothing / blurry masks $\rightarrow$ increase decoder capacity or resolution.
- (2) Rare class collapse $\rightarrow$ re-weighting, patch sampling, Dice/Tversky, focal-like terms.

Evaluation: Metrics and What They Diagnose

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

IoU and mIoU

For class k , with predicted set $\hat{Y}_k$ and ground-truth set Y_k :

$$IoU_k = \frac{|Y_k \cap \hat{Y}_k|}{|Y_k \cup \hat{Y}_k|}, \quad mIoU = \frac{1}{K} \sum_{k=1}^K IoU_k$$

Dice (F1 for masks)

$$Dice_k = \frac{2|Y_k \cap \hat{Y}_k|}{|Y_k| + |\hat{Y}_k|}$$

Dice is often more sensitive to **small structures** than IoU.

Qualitative checks (must-have)

Boundary quality, thin objects, small targets, failure under illumination / viewpoint change.

Analytics: Model Comparison and Selection Logic

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion


Comparative summary (semantic segmentation)

Family	Strength	Weakness	When to use	Notes
U-Net / Enc-Dec	strong locality, stable training	limited global context	small/medium data, medical	skip connections are key
DeepLab (ASPP)	multi-scale context	heavier, tuning needed	natural scenes, scale variance	good with strong backbone
Transformer (SegFormer)	global context, strong transfer	data hungry, compute	large data, domain shift	AdamW + warmup

Selection logic (engineering)

Real-time / embedded: lightweight encoder-decoder; highest accuracy: strong backbone + DeepLab/Transformer; small target / imbalance: Dice/Tversky + patch sampling.

IoU



Section 4. Object Detection

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Object Detection: Problem Setup

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Goal (set prediction)

Given image $x \in \mathbb{R}^{H \times W \times 3}$, predict a set of objects:

$$f_{\theta}(x) = \{(b_i, c_i, s_i)\}_{i=1}^N, \quad b_i = (x_i, y_i, w_i, h_i) \in \mathbb{R}^4,$$

$$c_i \in \{1, \dots, K\}, \quad s_i \in [0, 1]$$

Why it is hard

- Strong imbalance: **many negatives** vs few objects.
- Scale variation: small/large objects; occlusion; clutter.
- Joint **classification** + **localization**.



YOLO Labels: Matching Predictions to Ground Truth

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Training targets

Given GT objects $\{(b_j^*, c_j^*)\}_{j=1}^{N^*}$, assign each GT to one (or few) predictions. Modern YOLO variants use **dynamic assignment** (e.g., IoU + classification quality).

Simplified assignment view

- Define candidate predictions near the GT center and across scales.
- Compute quality $q(u)$ using IoU and/or classification logit.
- Select top- m candidates as positives; others negatives.

Practical note

Bad assignment $\Rightarrow$ unstable training (box regression fights classification). Most accuracy gains in modern detectors come from better **label assignment** + **loss design**.

YOLO Architecture

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

- YOLO follows a **Backbone-Neck-Head** design.
- **Backbone**: extracts hierarchical features from the input image.
- **Neck (FPN/PAN)**: fuses multi-scale features for small and large object detection.
- **Detection Head**: predicts bounding boxes, objectness scores, and class probabilities.
- Final predictions are decoded and filtered using **Non-Maximum Suppression (NMS)**.




Figure: Typical YOLO architecture (Backbone-Neck-Head pipeline) Redmon et al. [2016].

YOLO Loss: Box + Objectness + Class

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Composite objective

$$\mathcal{L}_{det} = \underbrace{\lambda_{box}}_{\text{Box}} \mathcal{L}_{box} + \underbrace{\lambda_{obj}}_{\text{Objectness}} \mathcal{L}_{obj} + \underbrace{\lambda_{cls}}_{\text{Class}} \mathcal{L}_{cls}$$

Box regression (IoU family)

Using IoU-based losses improves localization:

$$\mathcal{L}_{box} = 1 - \text{IoU}(b, b^*) \quad \text{or} \quad 1 - \text{GIoU}/\text{DIoU}/\text{CIoU}$$

Objectness + classification

- $\mathcal{L}_{obj}$: BCE on $p_{obj}(u)$ (positive for matched predictions, negative otherwise).
- $\mathcal{L}_{cls}$: BCE/CE on class distribution for positive predictions.
- Focal-like modulation is often used to down-weight easy negatives.

YOLO Inference: Decoding and NMS

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Decoding

Predicted offsets Δb are converted to image-space boxes b (scale/stride dependent). Compute score $s = p_{obj} \cdot \max_k p_{cls}(k)$ (typical).

Non-Maximum Suppression (NMS)

Given candidate boxes with scores, suppress duplicates:

- Sort boxes by score (descending).
- Keep highest-score box; remove boxes with $IoU > \tau$ to it.
- Repeat until list exhausted.

Engineering knobs

Confidence threshold τ_s and NMS IoU threshold τ control precision/recall trade-off.

YOLO: Training & Inference Algorithm

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Algorithm 2 YOLO Pipeline (Training + Inference)

Require: Dataset $\mathcal{D} = \{(x^{(n)}, \{(b_j^*, c_j^*)\})\}_{n=1}^M$, params θ , lr η

1: **for** epoch = 1, ..., E **do**

2: **for** mini-batch $\{(x, GT)\}$ **do**

3: $Z \leftarrow f_\theta(x)$

4: $\mathcal{A} \leftarrow \text{Assign}(Z, GT)$

5: $\mathcal{L} \leftarrow \lambda_{\text{box}} \mathcal{L}_{\text{box}}(Z, GT, \mathcal{A}) + \lambda_{\text{obj}} \mathcal{L}_{\text{obj}}(Z, \mathcal{A}) + \lambda_{\text{cls}} \mathcal{L}_{\text{cls}}(Z, GT, \mathcal{A})$

6: $\theta \leftarrow \theta - \eta \nabla_\theta \mathcal{L}$

7: **end for**

8: **end for**

9: **Inference:** $Z \leftarrow f_\theta(x)$; decode boxes and scores; keep boxes with $s > \tau_s$; apply NMS with IoU threshold τ .

▷ Dense predictions at multi-scales

▷ Match GT to predictions (positives)

Other Detectors (Comparison Table)

CV-2026

A.Kornaev,
K.Yakovlev

Where YOLO sits in the detector landscape

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Family	Paradigm	Key components	Post-proc	Typical use
YOLO	one-stage	FPN/PAN + dense head, dynamic assign	NMS	real-time, robotics
Faster R-CNN	two-stage	RPN proposals + ROI head	NMS	max accuracy (offline)
RetinaNet	one-stage	anchors + focal loss	NMS	strong baseline (imbalanced)
DETR-like	set prediction	transformer + Hungarian matching	no NMS	clean pipeline, research

Selection logic (practical)

Real-time / edge $\Rightarrow$ YOLO. Highest accuracy (when speed is secondary) $\Rightarrow$ Faster R-CNN / strong two-stage. Clean end-to-end (no NMS) and good scaling with data $\Rightarrow$ DETR-family.

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Section 5. Segment Anything Model (SAM)

What is SAM?

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Key Idea

The Segment Anything Model (SAM) is a foundation model for segmentation developed by Meta AI. It can segment any object in an image using different input prompts, making it highly versatile.

- Works with bounding boxes, clicks, or freeform masks as prompts.
- Trained on SA-1B dataset with over 1 billion masks.
- Enables zero-shot segmentation on unseen images.




Figure: SAM Model Overview Kirillov et al. [2023].

SAM Architecture

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

- **Image Encoder:** A large Vision Transformer (ViT) pre-trained using masked autoencoding.
- **Prompt Encoder:** Processes input prompts (points, boxes, masks).
- **Mask Decoder:** Predicts segmentation masks based on encoded image and prompts.




Figure: Segment Anything Model (SAM) Architecture.

Training and Dataset

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

- **SA-1B Dataset:** The largest segmentation dataset with 1 billion masks.
- **Semi-Automatic Annotation:** Human annotators refined AI-generated masks.
- **Generalization:** Trained to segment unseen objects and domains.




Figure: Segment Anything Model (SAM) Architecture.

How SAM Works?

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

- The user provides an input prompt (point, bounding box, or mask).
- SAM generates segmentation masks based on prompt and image encoding.
- Output can be refined interactively.




Figure: Segment Anything Model (SAM) Architecture.

Applications of SAM

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

- **Medical Imaging:** Tumor and organ segmentation.
- **Autonomous Driving:** Road scene segmentation.
- **AR/VR:** Object segmentation for augmented reality.
- **Robotics:** Object recognition and manipulation.




Figure: Medical Application of SAM.

Key Takeaways

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

- SAM is a powerful, zero-shot segmentation model.
- Works with multiple input prompts for flexible segmentation.
- Trained on the largest segmentation dataset, enabling generalization.
- Open-source and widely applicable across industries.

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

Section 6. Conclusion

Key Takeaways

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic
Segmentation

Object
Detection

Segment
Anything
Model (SAM)

Conclusion

- Segmentation and detection are essential for computer vision applications.
- Deep learning-based models have achieved state-of-the-art performance.
- Practical implementation involves dataset preparation, model training, and evaluation.

Bibliography

CV-2026

A.Kornaev,
K.Yakovlev

Outcomes

Introduction

Semantic Segmentation

Object Detection

Segment Anything Model (SAM)

Conclusion

Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C Berg, Wan-Yen Lo, et al. Segment anything. In **Proceedings of the IEEE/CVF International Conference on Computer Vision**, pages 4015–4026, 2023.

Simon J.D. Prince. **Understanding Deep Learning**. The MIT Press, 2023. URL <http://udlbook.com>.

Joseph Redmon et al. You only look once: Unified, real-time object detection. **CVPR**, 2016.

Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation, 2015. URL <https://arxiv.org/abs/1505.04597>.